

RTX 3090: four heads, $N_q=N_k=2048$, $d=d_v=128$
Rectangular inference microbenchmark; not end-to-end LLM throughput

